

3 General Discussion

3.1 Keys to Be Generated

This Recommendation addresses the generation of the cryptographic keys used in cryptography. Key generation includes the generation of a key using the output of a random bit generator (RBG), the derivation of a key from another key, the derivation of a key from a password, and key agreement performed by two entities using an **approved** key-agreement scheme. All keys **shall** be based directly or indirectly on the output of an **approved** RBG. For the purposes of this Recommendation, keys that are derived during a key-agreement transaction (see [SP 800-56A](#) and [SP 800-56B](#)), derived from another key using a key derivation function (see [SP 800-108](#)), or derived from a password for storage applications (see [SP 800-132](#)¹³ and [Section 6.5](#)) are considered to be indirectly generated from an RBG since an ancestor key¹⁴ or random value (e.g., the random value used to generate a key-agreement key pair) was obtained directly from the output of an **approved** RBG.

Two classes of cryptographic algorithms that require cryptographic keys have been **approved** for use by the Federal Government: asymmetric-key algorithms and symmetric-key algorithms. The generation of keys for these algorithm classes is discussed in [Sections 5](#) and [6](#), respectively.

3.2 Where Keys are Generated

Cryptographic keys **shall** be generated within [FIPS 140](#)-validated cryptographic modules. For explanatory purposes, consider the cryptographic module in which a key is generated to be the key-generating module. Any random value required by the key-generating module **shall** be generated within that module; that is, the RBG (or portion of the RBG¹⁵) that generates the random value **shall** be implemented within the FIPS 140 cryptographic module that generates the key. The generated keys **shall** be transported (when transportation is necessary) using secure channels and **shall** be used by their associated cryptographic algorithm within FIPS 140-validated cryptographic modules.

3.3 Supporting a Security Strength

A method (e.g., an RBG or a key and its associated cryptographic algorithm) *supports a given security strength* if the security strength provided by that method is equal to or greater than the security strength required for protecting the target data; the actual security strength provided can be higher than required.

Security strength supported by an RBG: A well-designed RBG supports a given security strength only if the amount of entropy (i.e., the randomness) available in the RBG is equal to or greater than that security strength. The support of a given security strength also requires a commensurate security strength for the confidentiality protection afforded to the entropy bits entered into the RBG (and other parameters determining the RBG's state); when used for the generation of keys and other

¹³ SP 800-132, *Recommendation for Password-Based Key Derivation, Part 1: Storage Applications*.

¹⁴ Ancestor key: A key that is used in the generation of another key. For example, an ancestor key for a key generated by a key derivation function would be the key-derivation key used by that key derivation function.

¹⁵ The RBG itself might be distributed (e.g., the entropy source may not co-reside with the algorithm that generates the (pseudo) random output).

secret values, a commensurate security strength is also required for the confidentiality and integrity protection that will be provided to the RBG output. For information regarding the security strength that can be supported by **approved** RBGs, see [SP 800-90A](#).¹⁶

Security strength supported by an algorithm: Discussions of cryptographic algorithms and the security strengths they can support given certain choices of parameters and/or key lengths are provided in [SP 800-57, Part 1](#).¹⁷ The security strength of a cryptographic algorithm that uses keys of a certain size (i.e., the key's length) is assessed under the assumption that those keys are generated using an **approved** process that outputs keys of the requisite size and type while providing an appropriate amount of min-entropy. It is assumed that this key-generation process is capable of supporting security strengths that are equal to or greater than the security strength assessed for the cryptographic algorithm. (Both the min-entropy and the security strength are measured in bits.)

Security strength supported by a key: The security strength that can be supported by a key depends on: 1) the algorithm with which it is used, 2) the size of the key (see [SP 800-57, Part 1](#)), 3) the process that generated the key (e.g., the security strength supported by the RBG that was used to generate the key), and 4) how the key was handled (e.g., the security strength available in the method used to transport the key). The use of such terms as “security strength supported by a key” or “key supports a security strength” assumes that these factors have been taken into consideration. For example, if an **approved** RBG that supports a security strength of 128 bits has been used to generate a 128-bit key, and if (immediately after generation) the key is used with AES-128 to encrypt target data, then the key may be said to support a security strength of 128 bits in that encryption operation (for as long as the key is kept secret). However, if the 128-bit AES key is generated using an RBG that supports a security strength of only 112 bits, then the key can only support a security strength of 112 bits even though its length is still 128 bits (i.e., the security strength of the key has been determined by the process used for its generation).

¹⁶ SP 800-90A, *Recommendation for Random Number Generation Using Deterministic Random Bit Generators*.

¹⁷ SP 800-57, Part 1, *Recommendation for Key Management: General*.

4 Using the Output of a Random Bit Generator

Random bit strings required for the generation of cryptographic keys **shall** be obtained from the output of an **approved** random bit generator (RBG); **approved** RBGs are specified in [SP 800-90](#).¹⁸ The RBG **shall** be instantiated at a security strength that supports the security strength required to protect the target data (i.e., the data that will be protected by the generated keys).

The output of an **approved** RBG can be used as specified in this section to obtain, for example, either a symmetric key or the random value needed to generate an asymmetric key pair.

Asymmetric key pairs require the use of an **approved** algorithm for their generation. Examples are those included in [FIPS 186](#) for generating DSA, ECDSA, and RSA keys. The generation of asymmetric key pairs from a random value is discussed in [Section 5](#).

Methods for the generation of symmetric keys are discussed in [Section 6](#).

When random bit strings are required for the generation of cryptographic keys, they are obtained as follows:

Let B be the random bit string to be acquired, for example, to use as a symmetric key (K), as input to an asymmetric-key-generation algorithm, or as part of the seed material for an RBG. Let $bLen$ be its desired length in bits. B **shall** be a bit string that is formed as follows:

$$B = U \oplus V,$$

where

- U is a bit string of $bLen$ bits that is obtained as the output of an **approved** RBG that is capable of supporting the security strength required by the algorithm and/or application using B (e.g., to protect the target data),
- V is a bit string of $bLen$ bits (which could be all zeroes), and
- The value of V is determined in a manner that is independent of the value of U (and vice versa).

The algorithm and/or application with which B will be used and the security strength that B is intended to support will determine the required bit length of B : $bLen$ **shall** meet the relying application's or algorithm's length requirement for a value of the bit string B to be used as intended in support of the targeted security strength. For example, if B is to be used as an AES key, then $bLen$ **shall** be an **approved** AES key length that supports the required security strength for protecting the target data. As another example, according to [FIPS 186](#), if B is to be used as a seed in the specified process of generating provably prime factors of an RSA modulus n , then $bLen$ **shall** be twice the security strength associated with (the bit length of) n .

Since there are no restrictions on the selection of V (other than its length and independence from U), a conservative approach necessitates an assumption that the process used to select U provides most (if not all) of the required entropy, which – when measured in bits – cannot exceed the length of U (i.e., $bLen$). Therefore, the **approved** RBG from which U is obtained **shall** be capable of

¹⁸ SP 800-90, *Recommendation for Random Number Generation*, consisting of SP 800-90A, SP 800-90B, and SP 800-90C.

providing the requisite entropy for B during the generation of U (i.e., at least $bLen$ bits of entropy are provided during the seeding of the RBG).

The independence requirement on U and V is interpreted in a computational and statistical sense: the computation of U does not depend on V , and the computation of V does not depend on U . Knowledge of the value of V (but not B) must provide no advantage to a party intent on gaining insight into an (as-yet-unknown) value of U . The value of V may be selected using a process that provides little entropy (e.g., V may be assigned a fixed, public value). Nevertheless, in cases where the value of V is intended to be kept secret, knowledge of the value of U (but not B) must yield no additional information concerning the value selected for V . Given that U is the output of an **approved** RBG, the following are examples of independently selected V values:

1. V is a constant (selected independently of the value of U). The value of V may be dependent on the use of B (e.g., if B is used as a key-derivation key, then V may be some value M , but if B is used as a key-wrapping key, then V may be some value N). Note that if V is a string of binary zeroes, then $B = U$ (i.e., the output of an **approved** RBG).
2. V is a key obtained using an **approved** key-derivation method from a key-derivation key and other input that is independent of U ; see [SP 800-108](#).
3. V is a key that was independently generated in another cryptographic module. V was protected using an **approved** key-wrapping algorithm or transported using an **approved** key-transport scheme during subsequent transport. Upon receipt, the protection on V is removed within the key-generating module that generated U before combining V with U .
4. V is produced by hashing another bit string (V') using an **approved** hash function and (if necessary) truncating the result to the appropriate length before combining it with U . That is, $V = T(H(V'), bLen)$ where $T(x, bLen)$ denotes the truncation of bit string x to its $bLen$ leftmost bits. The bit string V' may be selected using methods 1, 2, or 3 above.

5 Generation of Key Pairs for Asymmetric-Key Algorithms

Asymmetric-key algorithms (also known as public-key algorithms) require the use of asymmetric key pairs consisting of a private key and a corresponding public key. A key pair can be used for the generation and verification of digital signatures (see [Section 5.1](#)) or for key establishment (see [Section 5.2](#)). Each public/private key pair is associated with only one entity; this entity is known as the key-pair owner. Key pairs **shall** be generated by:

- The key-pair owner, or
- A Trusted Party that provides the key pair to the owner in a secure manner. The Trusted Party must be trusted by all parties that use the public key.

After key-pair generation, the key pair is retained and used by its owner. If the key pair was generated by a Trusted Party, both the owner and any relying party must trust that party not to use the private key of the key pair. The public key may be known by or provided to whomever needs to use it when interacting with the owner (see [Section 5.3](#)).

5.1 Key Pairs for Digital Signature Schemes

Digital signatures are generated on data to provide source authentication, identity authentication, assurance of data integrity, or support for signatory non-repudiation. Digital signatures are generated by a signer using a private key and verified by a receiver using a public key. The generation of key pairs for digital signature applications is addressed in [FIPS 186](#) for the DSA, RSA, and ECDSA digital signature algorithms.

Values of B , computed as shown in [Section 4](#), **shall** be used to provide all random bit strings used in key-pair generation (which includes the generation of “secret numbers” used as ephemeral private keys), as specified in [FIPS 186](#). The maximum security strength that can be supported by the resulting key pairs depends on a variety of size and parameter choices. Guidance on the size/parameter choices appropriate for supporting various security strengths can be found in [SP 800-57, Part 1](#).

For example, [SP 800-57, Part 1](#) states that an ECDSA key pair generated using an appropriate elliptic curve and a base point whose order is a 224-bit to 255-bit prime number can support (at most) a security strength of 112 bits. [FIPS 186](#) specifies that for such ECDSA key pairs, the random value used to determine a private key must be obtained using an RBG that supports a security strength of 112 bits. Using the method in [Section 4](#), a random value B that is to be used for the generation of the private key is determined by U (a value of a specified bit length obtained from an RBG that supports a security strength of at least 112 bits) and V (which could be zero). The value of B is then used to determine the private key from which the public ECDSA key is obtained, as specified in [FIPS 186](#).

5.2 Key Pairs for Key Establishment

Key establishment includes both key agreement and key transport. Key agreement is a method of key establishment in which the resultant secret keying material is a function of information contributed by all participants in the key-establishment process (usually only two participants) so that no party can predetermine the value of the keying material independent of any other party's contribution. For key-transport, one party (the sender) selects a value for the secret keying material

and then securely distributes that value to one or more other parties (the receiver(s)).

Approved methods for generating the asymmetric key pairs used by **approved** key-establishment schemes between two parties are specified in [SP 800-56A](#) (for schemes that use finite-field or elliptic-curve cryptography) and [SP 800-56B](#) (for schemes that use integer-factorization cryptography, such as RSA).

Values of B , computed as shown in [Section 4](#), **shall** be used to provide the random values¹⁹ needed to generate key pairs for the finite field or elliptic curve schemes in [SP 800-56A](#) or to generate key pairs for the integer-factorization schemes specified in [SP 800-56B](#). The maximum security strength that can be supported by the **approved** key-establishment schemes and the key sizes used by these schemes is provided in [SP 800-57, Part 1](#).

5.3 Distributing the Key Pairs

A general discussion of the distribution of asymmetric key pairs is provided in [SP 800-57, Part 1](#). Key pairs may either be static or ephemeral. Static key pairs are intended to be used multiple times; ephemeral keys are usually used only once.

The private key of a key pair **shall** be kept secret. It **shall** either be generated 1) within the key-pair owner's cryptographic module (i.e., the key-pair owner's key-generating module) or 2) within the cryptographic module of an entity trusted by the key-pair owner and any relying party not to misuse the private key or reveal it to other entities (i.e., the key pair is generated within the key-generating module of a Trusted Party and securely transferred to the key-pair owner's cryptographic module).

If a private key is ever output from a cryptographic module, the key **shall** be output and transferred in a form and manner that provides appropriate assurance²⁰ of its confidentiality and integrity (e.g., using manual methods and multi-party control procedures or automated key-transport methods). The protection **shall** provide appropriate assurance that only the key-pair owner and/or the party that generated the key pair will be able to determine the value of the plaintext private key (e.g., the confidentiality and integrity protection for the private key uses a cryptographic mechanism that is at least as strong as the (maximum) security strength that must be supported by the asymmetric-key algorithm that will use the private key).

The public key of a key pair may be made public. However, it **shall** be distributed and verified in a manner that assures its integrity and association with the key-pair owner (e.g., in the case of a static public key, this may be accomplished using an X.509 certificate that provides a level of cryptographic protection that is at least as strong as the security strength associated with the key pair).

5.4 Key Pair Replacement

Key pairs need to be replaced if the private key is compromised. Key pairs also need to be replaced occasionally to limit the amount of information that is protected by the key pair in case of a compromise of the private key (see Section 5.3 of [SP 800-57, Part 1](#)). Section 5.3.4 of SP 800-57,

¹⁹ Note that in Section 4, if V is all zeroes, then B (the random value) is the output of an RBG.

²⁰ The term "provide appropriate assurance" is used to allow various methods for the input and output of cryptographic keys to/from cryptographic modules that may be implemented at different security levels (see [FIPS 140](#) and Section 7.7 of the [FIPS 140 IG](#)).

Part 1 discusses the usage period for each key of the key pair for both digital signature and key-establishment key pairs.

When asymmetric key pairs need to be replaced, they **shall** be generated and distributed as specified in Sections [5.1](#), [5.2](#), or [5.3](#), as appropriate.

6 Generation of Keys for Symmetric-Key Algorithms

Symmetric-key algorithms use the same (secret) key to both apply cryptographic protection to information²¹ and to remove or verify the protection.²² Keys used with symmetric-key algorithms must be known by only the entities authorized to apply, remove, or verify the protection and are commonly known as secret keys. A secret key is often known by multiple entities that are said to share or own the secret key, although it is not uncommon for a key to be generated, owned, and used by a single entity (e.g., for secure storage). A secret key **shall** be generated by:

- One or more of the entities that will share the key, or
- A Trusted Party that provides the key to the intended sharing entities in a secure manner. The Trusted Party must be trusted (by all entities that will share the key) not to disclose the key to unauthorized parties or otherwise misuse the key (see [SP 800-71](#)²³).

A symmetric key K could be used, for example, to:

- Encrypt and decrypt data in an appropriate mode (e.g., using AES in the CTR mode, as specified in [FIPS 197](#) and [SP 800-38A](#));
- Generate Message Authentication Codes (e.g., using AES in the CMAC mode, as specified in [FIPS 197](#) and [SP 800-38B](#); HMAC, as specified in [FIPS 198](#); or KMAC, as specified in [SP 800-185](#)); or
- Derive additional keys using a key-derivation function specified in [SP 800-108](#), where K is the pre-shared (i.e., pre-existing) key that is used as the key-derivation key (e.g., K could be a value of B generated as specified in [Section 4](#)).

[Section 6.1](#) discusses the generation of symmetric keys that are obtained from the output of an RBG. [Section 6.2](#) discusses the derivation of symmetric keys. [Section 6.3](#) specifies **approved** techniques for combining a symmetric key with other symmetric keys and/or additional data.

At some point, a symmetric key needs to be replaced for a number of possible reasons (e.g., its cryptoperiod has been exceeded, or it has been compromised; see [SP 800-57 Part 1](#)). [Section 6.4](#) discusses key replacement.

6.1 The “Direct Generation” of Symmetric Keys

Symmetric keys that are to be directly generated from the output of an RBG **shall** be generated as specified in [Section 4](#), where B is used as the desired key K . The length of the key to be generated depends on the length requirement of the application or algorithm with which the key is used and the security strength to be supported. See [SP 800-57, Part 1](#) for discussions on key lengths and the (maximum) security strengths supported by symmetric-key algorithms and their keys.

²¹ For example, to transform plaintext data into ciphertext data using an encryption operation or compute a message authentication code (MAC).

²² For example, to remove the protection by transforming the ciphertext data back to the original plaintext data using a decryption operation or verify the protection by computing a message authentication code and comparing the newly computed MAC with a received MAC.

²³ SP 800-71, *Recommendation for Key Establishment Using Symmetric Block Ciphers*.

6.2 Derivation of Symmetric Keys

Symmetric keys are often obtained from the output of an **approved** key-derivation method (KDM), which is a cryptographic process specifically designed to transform secret input values into bit strings that can be parsed into cryptographic keys and/or other secret keying material.

Approved KDMs have been constructed from more basic cryptographic components, such as an **approved** hash function, as specified in [FIPS 180](#) or [FIPS 202](#); HMAC (using an **approved** hash function), as specified in [FIPS 198](#); AES-CMAC, as specified in [FIPS 197](#) and [SP 800-38B](#); or a KMAC variant, as specified in [SP 800-185](#).

Depending on the application and the KDM, the input to a KDM may include, for example, one or more of the following:

- A shared secret value produced during the execution of a key-agreement scheme;
- A cryptographic key (i.e., a key-derivation key (KDK));
- A password or passphrase;
- A salt value, which may be secret or non-secret, fixed, or randomly selected;
- A nonce (including RBG output) that may, for example, indicate the algorithm to be associated with the key (e.g., AES), the use of the key (e.g., email), or any other information that may be useful for associating a particular execution of the KDM with the key(s) to be derived.

Approved key-derivation methods can be divided into two categories:

- 1) The first category consists of one-step key-derivation methods, which are usually called key-derivation functions (KDFs). General-purpose KDFs are based on pseudorandom functions (PRFs) that use a KDK (and other input) to generate additional keys (see [SP 800-108](#)). Some special-purpose KDFs, which are employed only as components of key-agreement schemes, are used to obtain keying material from the shared secrets produced during the execution of such schemes (see [SP 800-56C](#) and [SP 800-135](#)); other special-purpose KDFs are to be used only for password-based protection of stored data and/or the keys that protect that data (see [SP 800-132](#)).
- 2) The second category consists of extraction-then-expansion key-derivation procedures that involve two steps:
 - a. Randomness extraction to obtain a single cryptographic key-derivation key. The extraction of a KDK from a shared secret produced during the execution of a key-agreement scheme is described in [SP 800-56C](#). The HMAC-based extraction of a symmetric key from the concatenation of pre-existing symmetric keys (and, perhaps, other data) is described in [Section 6.3](#) (along with other methods of combining preexisting keys to form a new key). The key resulting from a key-extraction process can be used as a KDK for key expansion.

- b. Key expansion to derive keying material from 1) the key-derivation key produced during randomness extraction and 2) other information, as specified in [SP 800-56C](#)²⁴ and [SP 800-108](#).

6.2.1 Symmetric Keys Generated Using Key-Agreement Schemes

When an **approved** key-agreement scheme is available within an entity's key-generating module, a symmetric key may be established with another entity that has the same capability. This process results in a symmetric key that is shared between the two entities participating in the key-agreement transaction.

[SP 800-56A](#) and [SP 800-56B](#) provide several methods for pairwise key agreement. Asymmetric key-agreement keys are used with a key-agreement primitive algorithm to generate a shared secret. The shared secret is provided to a key-derivation method to derive keying material. [SP 800-56C](#) specifies **approved** key-derivation methods for the key-agreement schemes in SP 800-56A and SP 800-56B.

The maximum security strength that can be supported by a key derived in this manner is dependent on: 1) the security strength supported by the asymmetric key pairs (as used during key establishment), 2) the key-derivation method used, 3) the length of the derived key, and 4) the algorithm with which the derived key will be used. See [SP 800-57, Part 1](#).

6.2.2 Symmetric Keys Derived from a Pre-existing Key

Symmetric keys are often derived using a key-derivation function (KDF) and a preexisting key known as a key-derivation key. For example, the preexisting key may have been:

- Generated from an **approved** RBG (see [Section 4](#)) and distributed as specified in [Section 6.4](#);
- Agreed upon using a key-agreement scheme (see [Section 6.2.1](#));
- Derived using a KDF and a (different) preexisting key as specified in [SP 800-108](#); or
- An **approved** function of multiple cryptographic keys (and, perhaps, other data) as described in [Section 6.3](#).

Approved methods for key derivation are provided in [SP 800-108](#), which specifies **approved** KDFs for deriving keys from a pre-shared (i.e., preexisting) key-derivation key. The KDFs are based on HMAC (as specified in [FIPS 198](#)) and CMAC (as specified in [SP 800-38B](#)).

If the derived keys need to be distributed to other entities, this may be accomplished as discussed in [Section 6.4](#).

In addition to the symmetric-key algorithm with which a derived key will be used, the security strength that can be supported by the derived key depends on the security strength supported by the key-derivation key and the KDF used (see [SP 800-57, Part 1](#) for the maximum security strength

²⁴ When the two-step key-derivation method is used by a key-establishment scheme.

that can be supported by HMAC and CMAC, and see [SP 800-107²⁵](#) for further discussions about the security strength of HMAC).

6.2.3 Symmetric Keys Derived from Passwords

In a number of popular applications, keys are generated from passwords. This is a questionable practice since passwords are usually selected using methods that provide very little entropy (i.e., randomness) and are, therefore, easily guessed. However, **approved** methods for deriving keys from passwords for storage applications²⁶ are provided in [SP 800-132](#). For these applications, users are strongly advised to select passwords using methods that provide a very large amount of entropy.

When a key is generated from a password, the entropy provided (and thus, the maximum security strength that can be supported by the generated key) **shall** be considered to be zero unless the password is generated using an **approved** RBG. In this case, the security strength that can be supported by the password (*password_strength*) is no greater than the minimum of the security strength supported by the RBG (*RBG_strength*) and the actual number of bits of RBG output (*RBG_outlen*) used in the password. That is, $password_strength \leq \min(RBG_strength, RBG_outlen)$.

6.3 Symmetric Keys Produced by Combining (Multiple) Keys and Other Data

When symmetric keys K_1, \dots, K_n are generated and/or established independently, they may be combined within a key-generating module to form a key K . Other items of data (D_1, \dots, D_m) can also be combined with the K_i to form K under the conditions specified below. Note that while the K_i values are required to be secret, the D_i values need not be kept secret.

The component symmetric keys (i.e., the K_i values) **shall** be generated and/or established independently (and subsequently protected as necessary) using **approved** methods²⁷ that support a security strength that is equal to or greater than the targeted security strength of the algorithm or application that will rely on the output key K . Each component key **shall** be kept secret and **shall not** be used for any purpose other than the computation of a specific symmetric key K (i.e., a given component key **shall not** be used to generate more than one key).

The independent generation/establishment of the component keys K_1, \dots, K_n is interpreted in a computational and a statistical sense; that is, the computation of any particular K_i value does not depend on any one or more of the other K_i values, and it is not feasible to use knowledge of any proper subset of the K_i values to obtain any information about the remaining K_i values.

When their use is permitted, D_1, \dots, D_m **shall** be generated or obtained using methods that ensure their independence from the values of the component keys K_1, \dots, K_n .

The required independence of the component keys from these other items of data is also interpreted in a computational and a statistical sense. This means that the computation of the K_i values does not depend on any of the D_j values, the computation of the D_j values does not depend on any of

²⁵ SP 800-107, *Recommendation for Applications Using Approved Hash Algorithms*.

²⁶ For example, inside a FIPS 140-validated cryptographic module.

²⁷ See Sections 4, 6.1, and 6.2.

the K_i values, and knowledge of the D_j values yields no information that can feasibly be used to gain insight into the K_i values. In cases where some (or all) of the D_j values are secret and the rest of the D_j values (if any) are public, “independence” also means that knowledge of the K_i values and public D_j values yields no information that can feasibly be used to gain insight into the secret D_j values.

Let K_1, \dots, K_n be the n component keys to be combined to form K . For each K_i (where $i = 1$ to n), let ss_Mi be the maximum security strength that can be supported by the combination of method(s) used to generate K_i and the method(s) used to protect it after generation (e.g., during key transport and/or storage). In particular, assume that an adversary that is capable of exerting an effort on the order of 2^{ss_Mi} “basic operations” of some sort will be able to compromise those methods and obtain the value of K_i .

The **approved** methods for combining the component keys and other data are:

1. Concatenating two or more keys, i.e.,

$$K = K_1 \parallel \dots \parallel K_n.$$

Notes:

- a. This method requires $n \geq 2$.
 - b. The sum of the bit lengths of the n component keys **shall** be equal to $kLen$, the required bit length for K .
 - c. The methods used to generate or establish the component keys **shall** be such that the sum of the min-entropies provided by those methods is equal to or greater than the min-entropy required for the resulting key K .
2. Exclusive-Oring one or more symmetric keys and possibly one or more other items of data, i.e.,

$$K = K_1 \oplus \dots \oplus K_n \oplus D_1 \oplus \dots \oplus D_m.$$

Notes:

- a. The length of each component key (K_i) and the length of each D_i **shall** be equal to $kLen$, the required bit length of K .
This method requires $m \geq 0$, $n \geq 1$ and $n + m \geq 2$.
 - If $m = 0$, then $D_1 \oplus \dots \oplus D_m$ is an all-zero bit string of bit length $kLen$.
 - If $m = 1$, then $D_1 \oplus \dots \oplus D_m$ is just D_1 .
 - If $n = 1$, then $K_1 \oplus \dots \oplus K_n$ is just K_1 and $D_1 \oplus \dots \oplus D_m$ **shall** be a non-zero bit string (in particular, m **shall** be at least 1 in this case).
 - b. The methods used to generate or establish the component keys **shall** be such that at least one of those methods provides min-entropy equal to or greater than the min-entropy required for the resulting key K .
3. A key-extraction process, i.e.,

$$K = T(\text{HMAC-hash}(\text{salt}, K_1 \parallel \dots \parallel K_n \parallel D_1 \parallel \dots \parallel D_m), kLen).$$

Notes:

- a. HMAC-*hash* **shall** be an implementation of HMAC (as specified in [FIPS 198](#), using an **approved** hash function *hash*) with a security strength that meets or exceeds the targeted security strength of the algorithm or application that will rely on the resulting key *K* (see [SP 800-57, part 1](#)).
- b. The *salt* is a secret or non-secret value with a length ≥ 0 that is used as the HMAC key. The *salt* must be known by all entities using this key-extraction process to obtain the same value of *K*.
- c. This method requires $n \geq 1$. If $n = 1$, then $K_1 \parallel \dots \parallel K_n$ is just K_1 .
- d. This method requires $m \geq 0$. If $m = 0$, then $D_1 \parallel \dots \parallel D_m$ is a null string; if $m = 1$, then $D_1 \parallel \dots \parallel D_m$ is just D_1 .
- e. *T* is the truncation function defined in [Section 2.3](#).
- f. The length of the output block of the hash function used with HMAC **shall** be at least *kLen* bits, the required bit length for *K*.
- g. The sum of the min-entropies provided by the methods used to generate or establish the component keys **shall** be equal to or greater than the min-entropy required for the output *K* and **should** be at least twice that amount of min-entropy.
- h. Alternative orderings are permitted when forming the concatenation of keys and data (including interleaving the keys and data), but the ordering must be known by all entities computing the value of *K*.
- i. The security strength of the key formed from combining multiple keys and data is subject to the considerations discussed in [Section 3.3](#).

6.4 Distributing Symmetric Keys

The symmetric key generated within a key-generating module often needs to be shared with one or more other entities that have their own cryptographic modules. The key may be distributed manually or using an **approved** key-transport or symmetric key-wrapping method (see [SP 800-56B](#), [SP 800-38F](#), and [SP 800-71](#)). See [SP 800-57, Part 1](#) for further discussion. The method used for key transport or key wrapping **shall** support the desired security strength needed to protect the target data (i.e., the data to be protected by the application or algorithm relying on the symmetric key). The requirements for the output of a key from a cryptographic module are discussed in [FIPS 140](#).

6.5 Replacement of Symmetric Keys

Sometimes, a symmetric key may need to be replaced. This may be due to a compromise of the key or the end of the key's cryptoperiod (see [SP 800-57, Part 1](#)). Replacement **shall** be accomplished through a rekeying process. Rekeying is the replacement of a key with a new key that is generated independent of the value of the old key (i.e., knowledge of the old key provides no knowledge of the value of the replaced key and vice versa).

When a compromised key is replaced, the new key **shall** be generated in a manner that provides assurance of its independence from the compromised key. The new key may be generated using any appropriate method in [Section 6](#) with the following restrictions:

1. The method used **shall** provide assurance that there is no feasibly detectable relationship between the new key and the compromised key. To that end, the new key **shall not** be derived or updated using the compromised key.
2. If the compromised key was generated in a manner that depended (in whole or in part) on a password (see Sections [6.2.3](#)), then that password **shall** be changed prior to the generation of any new key; in particular, the new key(s) **shall** be generated in a manner that is independent of the old password value.

If an uncompromised symmetric key is to be replaced, it **shall** be replaced using any method in [Section 6](#) that supports the required amount of security strength. However, if the key to be replaced was generated in a manner that depended (in whole or in part) on a password (see Sections [6.2.3](#)), that password **shall** be changed prior to the generation of the new key.